

# DATA SCIENCE 2 - BIG DATA

March 7, 2022


Data Science 2 - Big Data

Faculty of Mathematics and Physics

# BIG DATA

## INTRODUCTION

Volume of data is exploding:



Source: Data Age 2025, sponsored by Seagate with data from IDC Global DataSphere, Nov 2018

# BIG DATA

## INTRODUCTION



- ▶ The data volumes are exploding
- ▶ More data has been created in the past three years than in the entire previous history
- ▶ 175 zettabytes expected in 2025 = 175 trillion gigabytes

Prefix	Symbol	Base	Name
yotta	Y	$10^{24}$	septillion
zetta	Z	$10^{21}$	sextillion
exa	E	$10^{18}$	quintillion
peta	P	$10^{15}$	quadrillion
tera	T	$10^{12}$	trillion
giga	G	$10^9$	billion
mega	M	$10^6$	million
kilo	k	$10^3$	thousand

# HADOOP DISTRIBUTED FILE SYSTEM

## INTRODUCTION



- ▶ HDFS is a distributed, scalable, and portable file system written in Java
- ▶ A Hadoop instance is divided into HDFS and MapReduce:
  - ▶ HDFS is used for storing the data
  - ▶ MapReduce is used for processing data
- ▶ Hadoop YARN – responsible for managing computing resources in clusters and using them for scheduling users' applications
- ▶ Hadoop cluster has nominally a single namenode plus a cluster of datanodes,
- ▶ HDFS stores large files (typically in the range of gigabytes to terabytes).
- ▶ With the default data is stored on three nodes: two on the same rack, and one on a different rack.
- ▶ Data nodes can talk to each other to rebalance data, to move copies around, and to keep the replication of data high.

# HADOOP DISTRIBUTED FILE SYSTEM

## SERVICES



- ▶ Hadoop has multiple services:
  - ▶ Name Node (Master Node): contains the details of locations of the data and their replications
  - ▶ Data Node (Slave Node): stores data in it as blocks. Sends a Heartbeat message to the Name node every 3 seconds
  - ▶ Secondary Name Node: This is only to take care of the checkpoints of the file system metadata which is in the Name Node.
  - ▶ ResourceManager: Arbitrates resources among all applications in the system.
  - ▶ ApplicationMaster: negotiating resources from the ResourceManager and working with the NodeManager(s) to execute and monitor tasks
  - ▶ NodeManager: YARN's per-node agent: Keeping up-to-date with ResourceManager, overseeing individual tasks

# HADOOP DISTRIBUTED FILE SYSTEM

## MAPREDUCE: WORD COUNT



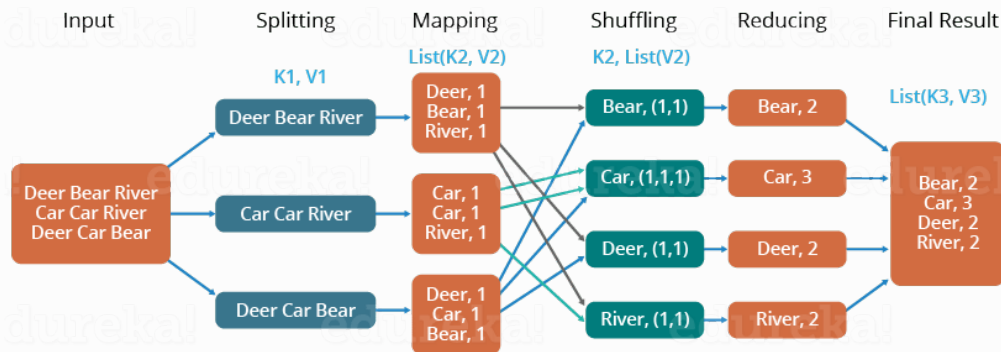
Most classical example is to count the number of occurrences of each word in the book:

- ▶ Let the prison warden order his guards to count the occurrence of words in his library.
- ▶ The guards, without hesitation, decided to involve prisoners in this task.
- ▶ Each prisoner can count occurrences in one book.

## BIG DATA

## MAPREDUCE: WORD COUNT

## The Overall MapReduce Word Count Process



# HADOOP DISTRIBUTED FILE SYSTEM

## MAPREDUCE



MapReduce is a programming paradigm model of using parallel, distributed algorithms to process or generate data sets. MapReduce is composed of two main functions:

- ▶ Map( $k,v$ ): Filters and sorts data.
- ▶ Reduce( $k,v$ ): Aggregates data according to keys ( $k$ ).

MapReduce is broken down into several steps:

- ▶ Record Reader
- ▶ Map
- ▶ Combiner (Optional)
- ▶ Partitioner
- ▶ Shuffle and Sort
- ▶ Reduce
- ▶ Output Format



# HADOOP DISTRIBUTED FILE SYSTEM

## MAPREDUCE: MAP



Record Reader translates an input into records of the form of a key-value pair  $(k_1, v_1)$ :

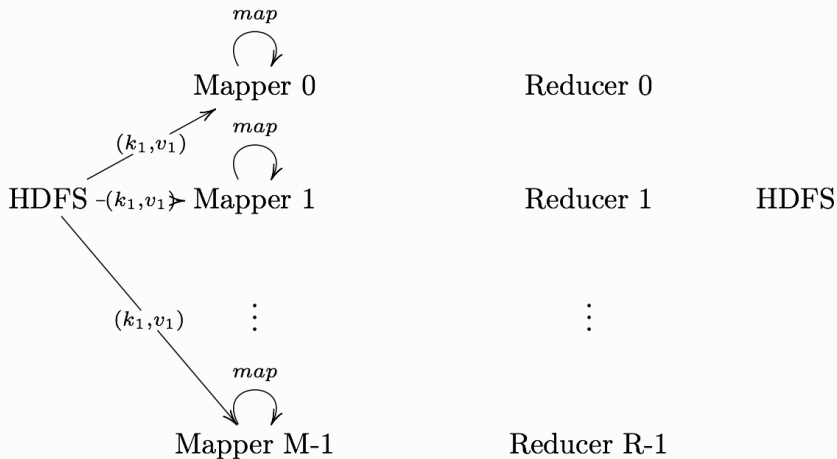
- ▶ These will be processed by the user-defined map function
- ▶ Key is positional information (the number of bytes from start of file) and the value is the chunk of data composing a single record.
- ▶ In hadoop, each map task's is an input split which is usually simply a HDFS block
- ▶ Hadoop tries scheduling map tasks on nodes where that block is stored (data locality)

Map is a user defined function outputting intermediate key-value pairs for the reducers:

- ▶  $\text{map}(k_1, v_1) \rightarrow \text{list}(k_2, v_2)$
- ▶ key  $(k_2)$ : Later, MapReduce will group and possibly aggregate data according to these keys, choosing the right keys is here is important for a good MapReduce job.
- ▶ value  $(v_2)$ : The data to be grouped according to it's keys.

# HADOOP DISTRIBUTED FILE SYSTEM

## MAPREDUCE: MAP



# HADOOP DISTRIBUTED FILE SYSTEM

## MAPREDUCE: COMBINER



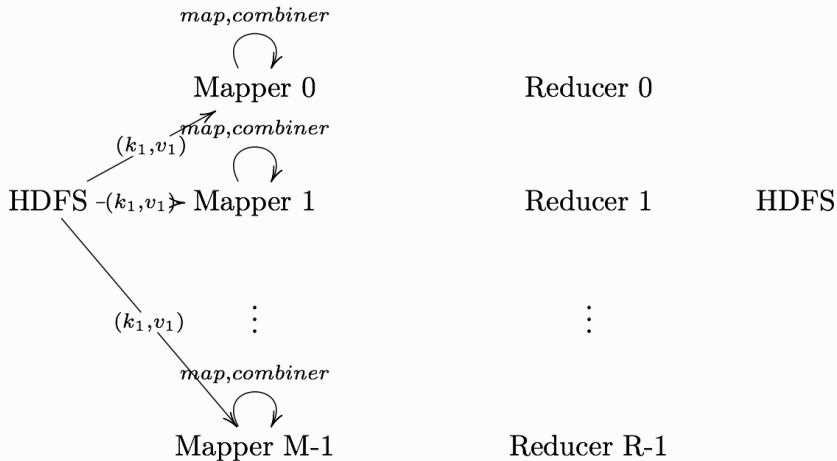
Combiner User defined function that aggregates data according to intermediate keys on a mapper:

- ▶ This can usually reduce the amount of data to be sent over the network increasing efficiency
- ▶ Combiner should be written with the idea that it is executed over most but not all map tasks.
- ▶ combiner:  $\text{list}(k2, v2) \rightarrow \text{list}(k2, v2)$

$$\left. \begin{array}{l} ("hello\ world", 1) \\ ("hello\ world", 1) \\ ("hello\ world", 1) \end{array} \right\} \xrightarrow{\text{combiner}} ("hello\ world", 3)$$

# HADOOP DISTRIBUTED FILE SYSTEM

## MAPREDUCE: COMBINER



# HADOOP DISTRIBUTED FILE SYSTEM

## MAPREDUCE: PARTITIONER



Partitioner sends intermediate key-value pairs  $(k, v)$  to reducer:


- ▶  $Reducer = \text{hash}(k) \pmod{R}$
- ▶ Should result in a roughly balanced load across the reducers while ensuring that all key-value pairs are grouped by their key on a single reducer.
- ▶ A balancer system is in place for the cases when the key-values are too unevenly distributed.
- ▶ In Hadoop, the intermediate keys  $(k_2, v_2)$  are written to the local harddrive and grouped by to which reducer they will be sent + the key itself.

Shuffle and Sort:

- ▶ On reducer node, sorts by key to help group equivalent keys

# HADOOP DISTRIBUTED FILE SYSTEM

## MAPREDUCE: SORT



# HADOOP DISTRIBUTED FILE SYSTEM

## MAPREDUCE: REDUCE



Reduce:

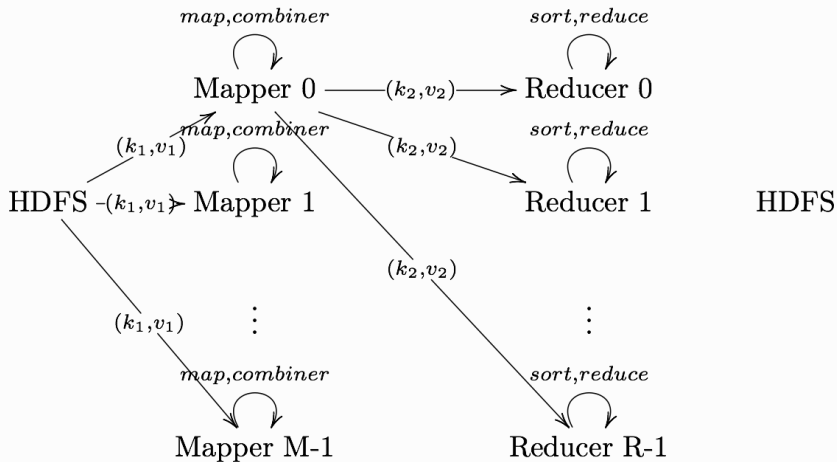
- ▶ User Defined Function that aggregates data ( $v$ ) according to keys ( $k$ ) to send key-value pairs to output:

Output Format:

- ▶ Translates final key-value pairs to file format (tab-separated by default).

# HADOOP DISTRIBUTED FILE SYSTEM

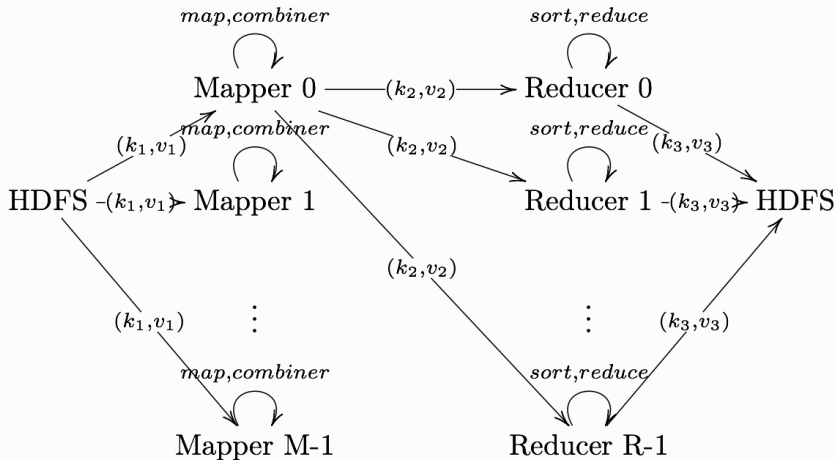
## MAPREDUCE: REDUCE





# HADOOP DISTRIBUTED FILE SYSTEM

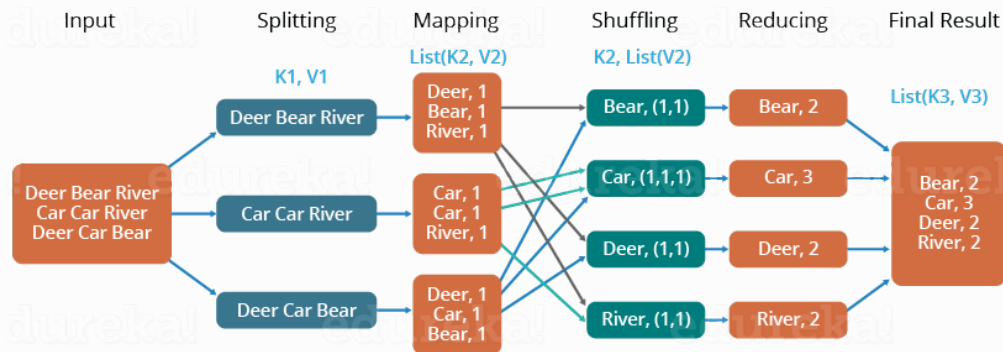
## MAPREDUCE: OUTPUT



# HADOOP DISTRIBUTED FILE SYSTEM

## MAPREDUCE: WORD COUNT

### The Overall MapReduce Word Count Process



# BIG DATA

## SPARK



- ▶ Resilient distributed dataset (RDD): read-only multiset of data items distributed over a cluster
- ▶ Tries to solve limitations in the MapReduce, which forces a particular linear dataflow structure
- ▶ Workflow is managed as a directed acyclic graph (DAG). Nodes represent RDDs while edges represent the operations on the RDDs.
- ▶ Supports both iterative algorithms (visit data set multiple times in a loop), and interactive/exploratory data analysis (repeated database-style querying of data)
- ▶ Requires a cluster manager and a distributed storage system:
  - ▶ native Spark cluster, Hadoop YARN, Kubernetes, etc.
  - ▶ HDFS, Amazon S3, Cassandra, etc.
- ▶ Spark SQL adds support for DataFrames and you can write SQL or Python scripts
- ▶ In Web UI, you can view the status of the jobs, cluster load and the decomposition over nodes
- ▶ Spark kernel available in Jupyter, you can transfer data between cluster and local machine

# BIG DATA

## SPARK DATAFRAME

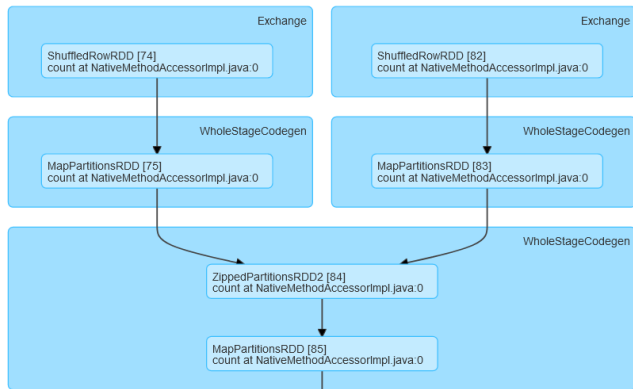


- ▶ Distributed table with schema (named and typed columns) instead of distributed collection (RDD)
- ▶ Dataframe uses all RDD mechanisms:
  - ▶ Partitioning
  - ▶ Transformations and actions
  - ▶ Non-modifiability
  - ▶ Lazy Computing
  - ▶ Caching a DataFrame in RAM and disk
- ▶ Pandas-like API
- ▶ Many formats for reading and saving data: csv, json, jdbc, hive, avro, incl. compact column-wise formats with indexes: parquet, orc
- ▶ Non-flat tables - column type support: struct, array, map
- ▶ User defined functions (UDF) with performance optimization tools written in python

# HADOOP DISTRIBUTED FILE SYSTEM

## SPARK UI

Stage 25




InMemoryTableScan

```

*(10) Project [ip#2859, created#2856, user_agent#2858, id#2852, click_id#2853, position#2854, section#2855, uuid#2857, bot_by_rule_2a#2951, bot_by_rule_2b#2974]
+- *(10) SortMergeJoin [ip#2859, created#2856], [ip#3006, created#3003], Inner
  :- *(6) Sort [ip#2859 ASC NULLS FIRST, created#2856 ASC NULLS FIRST], false, 0
  :- +- Exchange hashpartitioning(ip#2859, created#2856, 248)
  :- +- *(5) Project [created#2856, ip#2859, user_agent#2858, id#2852, click_id#2853, position#2854, section#2855, uuid#2857, bot_by_rule_2a#2951]
  :- +- *(5) SortMergeJoin [created#2856, ip#2859, user_agent#2858], [created#2986, ip#2989, user_agent#2988], Inner
  :- :- *(2) Sort [created#2856 ASC NULLS FIRST, ip#2859 ASC NULLS FIRST, user_agent#2858 ASC NULLS FIRST], false, 0
  :- :- +- Exchange hashpartitioning(created#2856, ip#2859, user_agent#2858, 248)
  :- :- +- *(1) Project [_source#2845 id AS id#2852, _id#2842 AS click_id#2853, _source#2845.position AS position#2854, _source#2... [86] [Cached]
count at NativeMethodAccessorImpl.java:0
  
```

**Thank you!**

TARAN



ADVISORY IN DATA & ANALYTICS